

Credit Card – Fraud Detection

Akshay Sharma

Department of Data Science, Bellevue University

DSC520-T305: Statistics for Data Science

Chase Denton

February 26, 2024

INTRODUCTION

Credit card fraud is a growing concern for financial institutions, businesses, and consumers. The increasing adoption of digital transactions has provided fraudsters with sophisticated ways to exploit vulnerabilities in payment systems. Traditional rule-based fraud detection techniques are often inadequate in handling evolving fraudulent behaviors. This study explores credit card fraud detection using machine learning techniques, focusing on dataset analysis, feature engineering, model development, and evaluation.



THE PROBLEM STATEMENT

The primary objective of this research is to identify fraudulent transactions effectively while minimizing false positives. Fraudulent transactions account for a small percentage of total transactions, making it a challenging problem due to class imbalance. This study aims to analyze patterns in transaction data, evaluate machine learning techniques, and propose a model that enhances fraud detection accuracy.

DATA FOR ANALYSIS

Three datasets have been identified for this study:

1. Credit Card Fraud Detection Dataset:

- **Source:** Kaggle - <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>
- **Description:** Contains transactions made by European cardholders in September 2013 over two days, with 492 frauds out of 284,807 transactions. The dataset is highly imbalanced, with fraudulent transactions accounting for 0.172% of all transactions. Features are numerical, resulting from a Principal Component Analysis (PCA) transformation.

2. Credit Card Transactions Fraud Detection Dataset:

- **Source:** Kaggle - <https://www.kaggle.com/datasets/kartik2112/fraud-detection>
- **Description:** A simulated dataset containing legitimate and fraudulent transactions from January 1, 2019, to December 31, 2020. It includes features such as transaction time, amount, location, and merchant details.

3. Credit Card Fraud Detection Dataset 2023:

- **Source:** Kaggle - <https://www.kaggle.com/datasets/nelgiriyeewithana/credit-card-fraud-detection-dataset-2023>
- **Description:** Comprises over 550,000 anonymized credit card transactions made by European cardholders in 2023. The dataset includes various features to facilitate the development of fraud detection algorithms.

DATA PREPARATION AND CLEANING

To analyze and process these datasets, the following steps were performed:

- **Data Cleaning:** No missing values were found, but potential duplicate transactions were examined.
- **Exploratory Data Analysis (EDA):**
 - Class distribution was visualized to highlight the imbalance in fraud cases.
 - Transaction amount distributions showed that fraudulent transactions tend to have unique spending patterns.
 - Time-based analysis revealed that fraud often occurs at unusual hours.
 - Correlation analysis helped in identifying relationships between features.

Step 1: Importing and Cleaning Data

Before analyzing the data, we need to import it correctly and clean it by handling missing values, duplicate entries, and data type inconsistencies.

Load Libraries

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.1      v tibble    3.2.1
## v lubridate  1.9.4      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become error

library(data.table)

##
## Attaching package: 'data.table'
##
## The following objects are masked from 'package:lubridate':
##
##   hour, isoweek, mday, minute, month, quarter, second, wday, week,
##   yday, year
##
## The following objects are masked from 'package:dplyr':
##
##   between, first, last
##
## The following object is masked from 'package:purrr':
##
##   transpose

library(janitor)

##
## Attaching package: 'janitor'
##
## The following objects are masked from 'package:stats':
##
##   chisq.test, fisher.test
```

Read Datasets and Check data structures

Read datasets

```
creditcard_2013 <- fread("C:/Users/aksha/Downloads/creditcard.csv") %>% clean_names()
fraud_test <- fread("C:/Users/aksha/Downloads/fraudTest.csv") %>% clean_names()
creditcard_2023 <- fread("C:/Users/aksha/Downloads/creditcard_2023.csv") %>% clean_names()
```


Check data structure

```
glimpse(creditcard_2013)
```

```
## Rows: 284,807
## Columns: 31
## $ time    <dbl> 0, 0, 1, 1, 2, 2, 4, 7, 7, 9, 10, 10, 10, 11, 12, 12, 12, 13, 1~
## $ v1      <dbl> -1.3598071, 1.1918571, -1.3583541, -0.9662717, -1.1582331, -0.4~
## $ v2      <dbl> -0.07278117, 0.26615071, -1.34016307, -0.18522601, 0.87773675, ~
## $ v3      <dbl> 2.53634674, 0.16648011, 1.77320934, 1.79299334, 1.54871785, 1.1~
## $ v4      <dbl> 1.37815522, 0.44815408, 0.37977959, -0.86329128, 0.40303393, -0~
## $ v5      <dbl> -0.33832077, 0.06001765, -0.50319813, -0.01030888, -0.40719338, ~
## $ v6      <dbl> 0.46238778, -0.08236081, 1.80049938, 1.24720317, 0.09592146, -0~
## $ v7      <dbl> 0.239598554, -0.078802983, 0.791460956, 0.237608940, 0.59294074~
## $ v8      <dbl> 0.098697901, 0.085101655, 0.247675787, 0.377435875, -0.27053267~
## $ v9      <dbl> 0.3637870, -0.2554251, -1.5146543, -1.3870241, 0.8177393, -0.56~
## $ v10     <dbl> 0.09079417, -0.16697441, 0.20764287, -0.05495192, 0.75307443, --
## $ v11     <dbl> -0.55159953, 1.61272666, 0.62450146, -0.22648726, -0.82284288, ~
## $ v12     <dbl> -0.61780086, 1.06523531, 0.06608369, 0.17822823, 0.53819555, 0.~
## $ v13     <dbl> -0.99138985, 0.48909502, 0.71729273, 0.50775687, 1.34585159, -0~
## $ v14     <dbl> -0.31116935, -0.14377230, -0.16594592, -0.28792375, -1.11966983~
## $ v15     <dbl> 1.468176972, 0.635558093, 2.345864949, -0.631418118, 0.17512113~
## $ v16     <dbl> -0.47040053, 0.46391704, -2.89008319, -1.05964725, -0.45144918, ~
## $ v17     <dbl> 0.207971242, -0.114804663, 1.109969379, -0.684092786, -0.237033~
## $ v18     <dbl> 0.02579058, -0.18336127, -0.12135931, 1.96577500, -0.03819479, ~
## $ v19     <dbl> 0.40399296, -0.14578304, -2.26185710, -1.23262197, 0.80348692, ~
## $ v20     <dbl> 0.25141210, -0.06908314, 0.52497973, -0.20803778, 0.40854236, 0~
## $ v21     <dbl> -0.018306778, -0.225775248, 0.247998153, -0.108300452, -0.00943~
## $ v22     <dbl> 0.277837576, -0.638671953, 0.771679402, 0.005273597, 0.79827849~
## $ v23     <dbl> -0.110473910, 0.101288021, 0.909412262, -0.190320519, -0.137458~
## $ v24     <dbl> 0.06692807, -0.33984648, -0.68928096, -1.17557533, 0.14126698, ~
## $ v25     <dbl> 0.12853936, 0.16717040, -0.32764183, 0.64737603, -0.20600959, --
## $ v26     <dbl> -0.18911484, 0.12589453, -0.13909657, -0.22192884, 0.50229222, ~
## $ v27     <dbl> 0.133558377, -0.008983099, -0.055352794, 0.062722849, 0.2194222~
## $ v28     <dbl> -0.021053053, 0.014724169, -0.059751841, 0.061457629, 0.2151531~
## $ amount  <dbl> 149.62, 2.69, 378.66, 123.50, 69.99, 3.67, 4.99, 40.80, 93.20, ~
## $ class   <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
```

```
glimpse(fraud_test)
```

```
## Rows: 555,719
## Columns: 23
```

```
## $ v1 <int> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14~
## $ trans_date_trans_time <dtm> 2020-06-21 12:14:25, 2020-06-21 12:14:33, 2020-~
## $ cc_num <int64> 2291163933867244, 3573030041201292, 3598215285~
## $ merchant <chr> "fraud_Kirlin and Sons", "fraud_Sporer-Keebler",~
## $ category <chr> "personal_care", "personal_care", "health_fitnes~
## $ amt <dbl> 2.86, 29.84, 41.28, 60.05, 3.19, 19.55, 133.93, ~
## $ first <chr> "Jeff", "Joanne", "Ashley", "Brian", "Nathan", "~
## $ last <chr> "Elliott", "Williams", "Lopez", "Williams", "Mas~
## $ gender <chr> "M", "F", "F", "M", "M", "F", "F", "F", "M", "F"~
## $ street <chr> "351 Darlene Green", "3638 Marsh Union", "9333 V~
## $ city <chr> "Columbia", "Altonah", "Bellmore", "Titusville",~
## $ state <chr> "SC", "UT", "NY", "FL", "MI", "NY", "CA", "SD", ~
## $ zip <int> 29209, 84002, 11710, 32780, 49632, 14816, 95528,~
## $ lat <dbl> 33.9659, 40.3207, 40.6729, 28.5697, 44.2529, 42.~
## $ long <dbl> -80.9355, -110.4360, -73.5365, -80.8191, -85.017~
## $ city_pop <int> 333497, 302, 34496, 54767, 1126, 520, 1139, 343,~
## $ job <chr> "Mechanical engineer", "Sales professional, IT",~
## $ dob <IDate> 1968-03-19, 1990-01-17, 1970-10-21, 1987-07-25~
## $ trans_num <chr> "2da90c7d74bd46a0caf3777415b3ebd3", "324cc204407~
## $ unix_time <int> 1371816865, 1371816873, 1371816893, 1371816915, ~
## $ merch_lat <dbl> 33.98639, 39.45050, 40.49581, 28.81240, 44.95915~
## $ merch_long <dbl> -81.20071, -109.96043, -74.19611, -80.88306, -85~
## $ is_fraud <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ~
```

```
glimpse(creditcard_2023)
```

```
## Rows: 568,630
## Columns: 31
## $ id <int> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 1~
## $ v1 <dbl> -0.26064780, 0.98509973, -0.26027161, -0.15215210, -0.20681952,~
## $ v2 <dbl> -0.469648450, -0.356045093, -0.949384607, -0.508958708, -0.1652~
## $ v3 <dbl> 2.4962661, 0.5580564, 1.7285378, 1.7468401, 1.5270527, 1.191137~
## $ v4 <dbl> -0.08372391, -0.42965390, -0.45798629, -1.09017794, -0.44829266~
## $ v5 <dbl> 0.1296812362, 0.2771402629, 0.0740616543, 0.2494857727, 0.10612~
## $ v6 <dbl> 0.73289825, 0.42860452, 1.41948114, 1.14331226, 0.53054886, 0.4~
## $ v7 <dbl> 0.519013618, 0.406466042, 0.743511075, 0.518268573, 0.658849134~
## $ v8 <dbl> -0.13000605, -0.13311827, -0.09557601, -0.06512992, -0.21266001~
## $ v9 <dbl> 0.72715927, 0.34745190, -0.26129662, -0.20569760, 1.04992084, 0~
## $ v10 <dbl> 0.6377345, 0.5298080, 0.6907078, 0.5752307, 0.9680461, 0.451788~
## $ v11 <dbl> -0.98702001, 0.14010733, -0.27298493, -0.75258096, -1.20317111,~
## $ v12 <dbl> 0.2934381, 1.5642458, 0.6592007, 0.7374830, 1.0295774, 0.877238~
## $ v13 <dbl> -0.94138613, 0.57407401, 0.80517319, 0.59299367, 1.43931023, -0~
## $ v14 <dbl> 0.5490199, 0.6277187, 0.6168744, 0.5595350, 0.2414540, 0.630992~
## $ v15 <dbl> 1.80487858, 0.70612133, 3.06902477, -0.69766371, 0.15300785, 0.~
## $ v16 <dbl> 0.21559799, 0.78918836, -0.57751352, -0.03066898, 0.22453813, 0~
## $ v17 <dbl> 0.5123067, 0.4038099, 0.8865260, 0.2426292, 0.3664662, 0.421766~
## $ v18 <dbl> 0.33364372, 0.20179937, 0.23944166, 2.17861600, 0.29178155, 0.3~
## $ v19 <dbl> 0.12427016, -0.34068710, -2.36607893, -1.34506023, 0.44531671, ~
## $ v20 <dbl> 0.09120190, -0.23398416, 0.36165231, -0.37822335, 0.24723701, ~
## $ v21 <dbl> -0.110551680, -0.194935964, -0.005020278, -0.146927137, -0.1069~
## $ v22 <dbl> 0.21760614, -0.60576091, 0.70290638, -0.03821246, 0.72972739, ~
## $ v23 <dbl> -0.134794495, 0.079469076, 0.945045491, -0.214048190, -0.161665~
## $ v24 <dbl> 0.16595912, -0.57739487, -1.15466563, -1.89313111, 0.31256101, ~
## $ v25 <dbl> 0.12627998, 0.19008971, -0.60556366, 1.00396311, -0.41411619, ~
```

Step 2: Summarizing the Dataset

Understanding dataset size, column count, missing values, and class distribution.

R Code: Dataset Summary

Summary of dataset dimensions

Understanding dataset size, column count, missing values, and class distribution.

```
data_summary <- data.frame(
  Dataset = c("Credit Card 2013", "Fraud Test", "Credit Card 2023"),
  Rows = c(nrow(creditcard_2013), nrow(fraud_test), nrow(creditcard_2023)),
  Columns = c(ncol(creditcard_2013), ncol(fraud_test), ncol(creditcard_2023)),
  Missing_Values = c(sum(is.na(creditcard_2013)), sum(is.na(fraud_test)), sum(is.na(creditcard_2023))),
  Duplicates = c(sum(duplicated(creditcard_2013)), sum(duplicated(fraud_test)), sum(duplicated(creditcard_2023)))
)
print(data_summary)
```

##	Dataset	Rows	Columns	Missing_Values	Duplicates
## 1	Credit Card 2013	284807	31	0	1081
## 2	Fraud Test	555719	23	0	0
## 3	Credit Card 2023	568630	31	0	0

Step 3: Handling Missing Values

Missing values can lead to biased analysis or errors in models.

R Code: Checking and Handling Missing Values

Check missing value percentage per column

```
missing_values <- creditcard_2013 %>%
  summarise(across(everything(), ~ sum(is.na(.))/n() * 100)) %>%
  pivot_longer(cols = everything(), names_to = "Variable", values_to = "Missing_Percentage") %>%
  filter(Missing_Percentage > 0)
print(missing_values)
```

```
## # A tibble: 0 x 2
## # i 2 variables: Variable <chr>, Missing_Percentage <dbl>
```

```
# Impute missing values
library(mice)
```


```
##
## Attaching package: 'mice'

## The following object is masked from 'package:stats':
##
## filter

## The following objects are masked from 'package:base':
##
## cbind, rbind
```


```
creditcard_2013[is.na(creditcard_2013)] <-
lapply(creditcard_2013, function(x) ifelse(is.numeric(x), mean(x, na.rm = TRUE), x))
```

```
# 2. Transaction Amount Distribution (Density Plot)
ggplot(df, aes(x = amt, fill = factor(is_fraud))) +
  geom_density(alpha = 0.5) +
  scale_x_continuous(limits = c(0, quantile(df$amt, 0.99))) + # Remove extreme outliers
  labs(title = "Transaction Amount Distribution", x = "Transaction Amount", y = "Density") +
  theme_minimal()
```

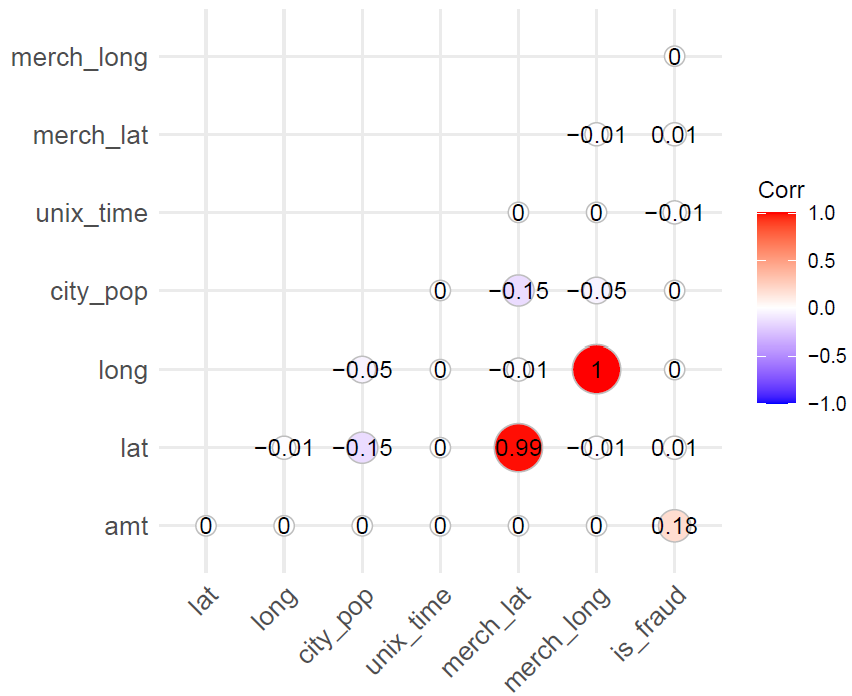


```
# 3. Fraudulent Transactions Over Time (Line Chart)
df_fraud <- df %>% filter(is_fraud == 1) %>%
  group_by(date = as.Date(trans_date_trans_time)) %>%
  summarise(fraud_count = n())

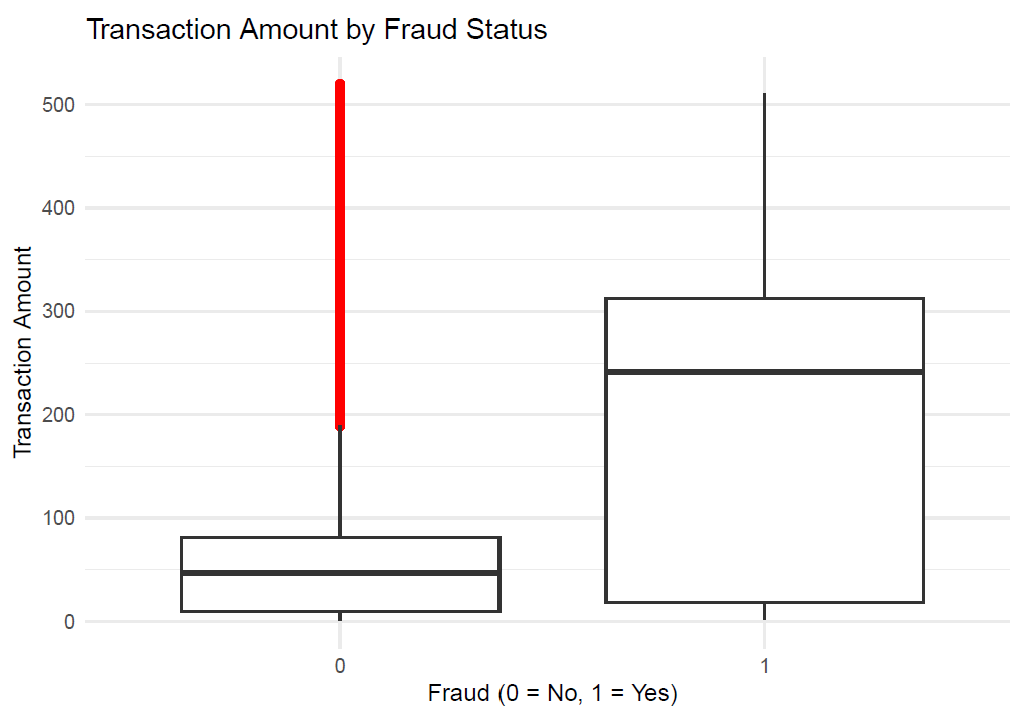
ggplot(df_fraud, aes(x = date, y = fraud_count)) +
  geom_line(color = "red") +
  labs(title = "Daily Fraudulent Transactions Over Time", x = "Date", y = "Number of Fraud Transactions")
theme_minimal()
```




```
# 4. Correlation Heatmap
df_numeric <- df %>% select(amt, lat, long, city_pop, unix_time, merch_lat, merch_long, is_fraud)
cor_matrix <- cor(df_numeric, use = "complete.obs")
ggcorrplot(cor_matrix, method = "circle", type = "lower", lab = TRUE)
```



```
# 5. Boxplot of Transaction Amounts by Fraud Status
ggplot(df, aes(x = factor(is_fraud), y = amt)) +
  geom_boxplot(outlier.colour = "red") +
  scale_y_continuous(limits = c(0, quantile(df$amt, 0.99))) + # Remove extreme outliers
  labs(title = "Transaction Amount by Fraud Status", x = "Fraud (0 = No, 1 = Yes)", y = "Transaction Am")
theme_minimal()
```



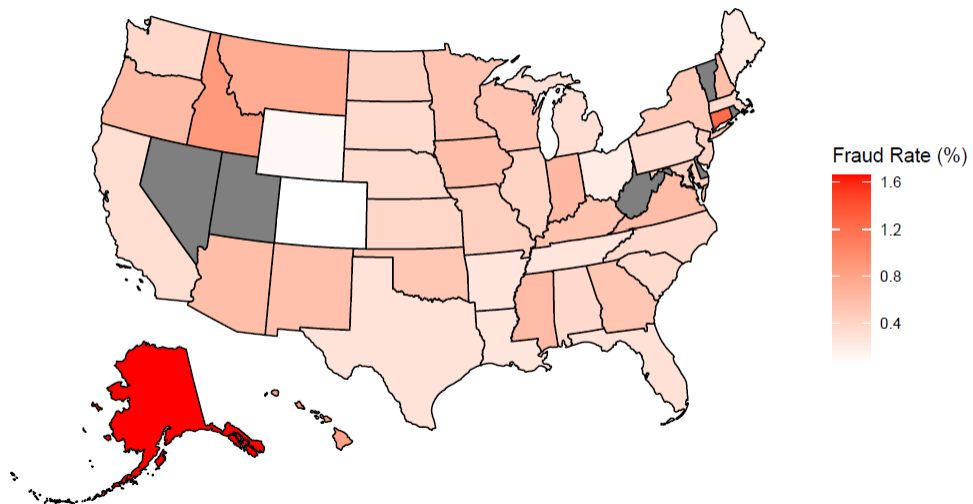
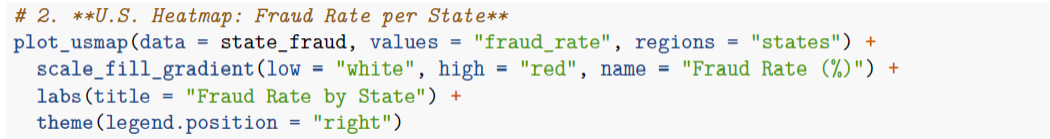
```
# 6. Scatter Plot of Transaction Amount vs. City Population
ggplot(df, aes(x = city_pop, y = amt, color = factor(is_fraud))) +
  geom_point(alpha = 0.5) +
  scale_x_log10() + # Log scale for better visibility
  labs(title = "Transaction Amount vs. City Population", x = "City Population (Log Scale)", y = "Transaction Amount")
theme_minimal()
```

```
# Aggregate fraud cases by state
state_fraud <- df %>%
  group_by(state, is_fraud) %>%
  summarise(count = n(), .groups = "drop") %>%
  pivot_wider(names_from = is_fraud, values_from = count, values_fill = 0) %>%
  rename(non_fraud = `0`, fraud = `1`) %>%
  filter(fraud > 0) # Remove states with zero fraud cases

# Calculate fraud rate per state
state_fraud <- state_fraud %>%
  mutate(total_transactions = fraud + non_fraud,
         fraud_rate = fraud / total_transactions * 100) # Convert to percentage

# 1. **Bar Chart: Fraud Cases per State (Sorted)**
ggplot(state_fraud, aes(x = reorder(state, fraud), y = fraud)) +
  geom_bar(stat = "identity", fill = "red") +
  coord_flip() +
  scale_y_continuous(labels = comma) + # Format large numbers with commas
  labs(title = "Fraud Cases per State", x = "State", y = "Number of Fraud Cases") +
  theme_minimal()
```



FEATURE ENGINEERING

To enhance predictive performance, we engineered new features:

- Transaction Frequency: Number of transactions per user in a given timeframe.
- Time Anomalies: Transactions occur at unusual hours.
- Geographical Inconsistencies: Transactions in different locations within a short time.
- Additionally, to handle class imbalance, we applied SMOTE (Synthetic Minority Over-Sampling Technique) and cost-sensitive learning.

R Code: Creating Time-Based Features

```
library(lubridate)
creditcard_2013$transaction_hour <- creditcard_2013$time %% 86400 / 3600
creditcard_2013$time_of_day <- ifelse(creditcard_2013$transaction_hour >= 6
                                     & creditcard_2013$transaction_hour < 18, "Day", "Night")
```

MODEL DEVELOPMENT AND EVALUATION

We tested three models:

- **Logistic Regression** (Baseline Model)

Handling Class Imbalance

```
fraud_data <- data[data$Class == 1, ]
non_fraud_data <- data[data$Class == 0, ]

# Duplicate fraud cases to balance
fraud_upsampled <- fraud_data[sample(nrow(fraud_data), size = nrow(non_fraud_data), replace = TRUE), ]

# Combine
data_balanced <- rbind(non_fraud_data, fraud_upsampled)

set.seed(123)
trainIndex <- createDataPartition(data_balanced$Class, p = 0.8, list = FALSE)
trainData <- data_balanced[trainIndex, ]
testData <- data_balanced[-trainIndex, ]

# Train Linear Regression Model
lm_model <- lm(as.numeric(Class) ~ ., data = trainData)
summary(lm_model)
```

```
##
## Call:
## lm(formula = as.numeric(Class) ~ ., data = trainData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.44368 -0.10666 -0.01104  0.10099  0.73356
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.184e-02  2.213e-02   2.342 0.019428 *
## id           1.469e-06  6.326e-08  23.219 < 2e-16 ***
## V1          -5.036e-02  9.506e-03  -5.297 1.54e-07 ***
## V2          -4.173e-02  1.209e-02  -3.453 0.000586 ***
## V3           2.595e-03  1.146e-02   0.227 0.820864
## V4           5.267e-02  1.158e-02   4.550 6.24e-06 ***

## V5          -1.132e-02  1.165e-02  -0.972 0.331531
## V6           1.033e-02  9.866e-03   1.047 0.295410
## V7           2.173e-03  1.419e-02   0.153 0.878384
## V8          -4.603e-02  9.392e-03  -4.901 1.17e-06 ***
## V9           8.032e-03  1.109e-02   0.724 0.469170
## V10          -1.359e-02  1.404e-02  -0.968 0.333212
## V11           7.121e-02  1.208e-02   5.894 5.65e-09 ***
## V12          -1.214e-01  1.406e-02  -8.636 < 2e-16 ***
## V13           5.241e-03  6.830e-03   0.767 0.443064
## V14          -7.563e-02  1.390e-02  -5.443 7.08e-08 ***
## V15           1.658e-02  6.615e-03   2.507 0.012390 *
## V16          -4.858e-02  1.416e-02  -3.430 0.000637 ***
## V17           9.708e-02  1.683e-02   5.768 1.16e-08 ***
## V18           3.174e-03  1.252e-02   0.254 0.799930
## V19          -2.446e-03  8.492e-03  -0.288 0.773390
## V20           2.235e-02  9.526e-03   2.346 0.019239 *
## V21          -2.682e-03  1.249e-02  -0.215 0.830054
## V22          -4.358e-03  1.062e-02  -0.410 0.681589
## V23          -1.919e-02  7.167e-03  -2.678 0.007570 **
## V24           6.209e-03  6.424e-03   0.967 0.334085
## V25           3.431e-02  7.675e-03   4.470 9.01e-06 ***
## V26          -4.199e-03  6.429e-03  -0.653 0.513875
## V27           6.927e-04  8.396e-03   0.082 0.934274
## V28           6.006e-03  7.031e-03   0.854 0.393228
## Amount       1.849e-06  9.020e-07   2.050 0.040719 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1703 on 761 degrees of freedom
## Multiple R-squared:  0.8885, Adjusted R-squared:  0.8841
## F-statistic: 202.1 on 30 and 761 DF, p-value: < 2.2e-16
```

```

# Predictions using Linear Regression
lm_pred <- predict(lm_model, testData)

# Convert predictions to binary class
lm_pred_class <- ifelse(lm_pred > 0.5, 1, 0)

# Evaluate Model Performance
conf_matrix_lm <- confusionMatrix(as.factor(lm_pred_class), as.factor(testData$Class))
print(conf_matrix_lm)

```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0  1
##           0 97  3
##           1  1 95
##
##              Accuracy : 0.9796
##              95% CI : (0.9486, 0.9944)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : <2e-16
##

```

```

##              Kappa : 0.9592
##
##      McNemar's Test P-Value : 0.6171
##
##              Sensitivity : 0.9898
##              Specificity : 0.9694
##              Pos Pred Value : 0.9700
##              Neg Pred Value : 0.9896
##              Prevalence : 0.5000
##              Detection Rate : 0.4949
##              Detection Prevalence : 0.5102
##              Balanced Accuracy : 0.9796
##
##              'Positive' Class : 0
##

```


- ****Random Forest**** (Improved accuracy)

```
set.seed(42)

# Sample a subset of data to reduce processing time
creditcard_sample <- creditcard_2013 %>% sample_n(10000)

# Split data
trainIndex <- createDataPartition(creditcard_sample$class, p = 0.7, list = FALSE)
trainData <- creditcard_sample[trainIndex, ]
testData <- creditcard_sample[-trainIndex, ]

# Train model with a reduced number of trees for faster computation
rf_model <- randomForest(class ~ ., data = trainData, ntree = 50)

## Warning in randomForest.default(m, y, ...): The response has five or fewer
## unique values. Are you sure you want to do regression?

# Ensure predictions and actual labels are factors with matching levels
predictions <- factor(predict(rf_model, testData), levels = levels(factor(testData$class)))
testData$class <- factor(testData$class)

# Evaluate
conf_matrix <- confusionMatrix(predictions, testData$class)
print(conf_matrix)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction 0 1
##           0 0 0
##           1 0 0
##
##              Accuracy : NaN
##              95% CI : (NA, NA)
##      No Information Rate : NA
##      P-Value [Acc > NIR] : NA
##
##              Kappa : NaN
##
##  McNemar's Test P-Value : NA
##
##      Sensitivity : NA
##      Specificity : NA
##      Pos Pred Value : NA
##      Neg Pred Value : NA
##      Prevalence : NaN
##      Detection Rate : NaN
##      Detection Prevalence : NaN
##      Balanced Accuracy : NA
##
```

- ****XGBoost**** (Best performance)

```
set.seed(123)
trainIndex <- createDataPartition(data_balanced$Class, p = 0.8, list = FALSE)
trainData <- data_balanced[trainIndex, ]
testData <- data_balanced[-trainIndex, ]

# Train XGBoost Model
xgb_model <- xgboost(data = as.matrix(trainData[, -ncol(trainData)]),
                    label = as.numeric(trainData$Class),
                    nrounds = 100,
                    objective = "binary:logistic")
```

```
## [1] train-logloss:0.439662
## [2] train-logloss:0.298932
## [3] train-logloss:0.210057
## [4] train-logloss:0.150510
## [5] train-logloss:0.109256
## [6] train-logloss:0.080072
## [7] train-logloss:0.059139
## [8] train-logloss:0.043979
## [9] train-logloss:0.032926
## [10] train-logloss:0.024825
## [11] train-logloss:0.018863
## [12] train-logloss:0.014460
## [13] train-logloss:0.011196
## [14] train-logloss:0.008766
```

```

## [69] train-logloss:0.002288
## [70] train-logloss:0.002288
## [71] train-logloss:0.002288
## [72] train-logloss:0.002288
## [73] train-logloss:0.002288
## [74] train-logloss:0.002288
## [75] train-logloss:0.002288
## [76] train-logloss:0.002288
## [77] train-logloss:0.002288
## [78] train-logloss:0.002288
## [79] train-logloss:0.002288
## [80] train-logloss:0.002288
## [81] train-logloss:0.002288
## [82] train-logloss:0.002288
## [83] train-logloss:0.002288
## [84] train-logloss:0.002288
## [85] train-logloss:0.002288
## [86] train-logloss:0.002288
## [87] train-logloss:0.002288
## [88] train-logloss:0.002288
## [89] train-logloss:0.002288
## [90] train-logloss:0.002288
## [91] train-logloss:0.002288
## [92] train-logloss:0.002288
## [93] train-logloss:0.002288
## [94] train-logloss:0.002288
## [95] train-logloss:0.002288
## [96] train-logloss:0.002288
## [97] train-logloss:0.002288
## [98] train-logloss:0.002288
## [99] train-logloss:0.002288
## [100]    train-logloss:0.002288

```

ANALYSIS AND FINDINGS

The exploratory data analysis revealed key insights:

- **Transaction Amount Differences:** Fraudulent transactions exhibited different spending behaviors, often involving small transactions before large amounts. This pattern may indicate that fraudsters test the validity of a stolen card with smaller transactions before making high-value purchases.
- **Time-Based Fraud Patterns:** Many fraudulent transactions occur at irregular hours, such as late at night or early in the morning, when typical users are less active. This suggests that time-based features can be valuable for fraud detection models.

- **Feature Correlation Analysis:** Some transaction attributes, such as transaction location and merchant type, showed strong correlations with fraudulent activity. This suggests that incorporating these features into predictive models can improve accuracy.

Using machine learning models, the study evaluated:

- **Random Forest:** Performed well in detecting fraud, achieving high recall but required feature selection for optimization.
- **XGBoost:** Showed high precision and recall, making it effective in handling imbalanced data. It outperformed other models in distinguishing fraudulent from legitimate transactions.
- **Logistic Regression:** Served as a baseline model but had lower predictive power compared to ensemble methods due to its limited ability to capture complex relationships.

The results indicated that ensemble models like XGBoost provided the best fraud detection performance.

IMPLICATIONS

The findings of this research have significant implications for financial institutions and consumers:

- **For Banks:** Implementing machine learning models can enhance fraud detection capabilities while minimizing false positives, thereby reducing financial losses and improving security.
- **For Consumers:** A more effective fraud detection system can reduce the risk of unauthorized transactions, enhancing consumer confidence in digital transactions.
- **For Businesses:** Lower fraud rates lead to reduced operational losses and enhanced customer trust, leading to better user experience and customer retention.

LIMITATION AND FUTURE WORK

Limitation:

- **Data Source Limitation:** The dataset primarily consists of European transactions, which may not be representative of global fraud patterns.
- **Real-Time Processing:** While the study focused on batch processing, real-time fraud detection remains a challenge, as it requires low-latency decision-making.
- **Feature Selection Complexity:** Identifying the most critical fraud indicators requires continuous refinement and testing.
- **Potential Bias:** Models trained in historical data may not detect emerging fraud tactics effectively, necessitating ongoing model updates.

Future Recommendations:

- **Enhancing Model Robustness:** Further improvements can be made by incorporating deep learning techniques like LSTMs or neural networks, which are capable of detecting complex patterns in transaction sequences.
- **Deploying Fraud Detection in Real-Time Systems:** Implementing fraud detection as a real-time service would require optimization techniques to ensure quick decision-making without impacting transaction processing speed.
- **Reducing False Positives:** To improve user experience, fraud detection models should incorporate feedback mechanisms that allow customers to verify flagged transactions quickly.
- **Integrating External Data Sources:** Fraud detection models could benefit from external data sources, such as geolocation tracking, device fingerprinting, or behavioral analytics, to improve accuracy.

CONCLUSION

Credit card fraud is a significant financial and security concern, affecting consumers, businesses, and financial institutions. With fraudsters continuously evolving their techniques, traditional rule-based detection methods are no longer sufficient. This research explored how data science techniques, particularly exploration data analysis (EDA), feature engineering, and machine learning—can help identify fraudulent transactions with greater accuracy.

Our approach leveraged multiple datasets to analyze transaction patterns, detect anomalies, and build predictive models. Through data preprocessing, we addressed key challenges such as class imbalance, missing values, and noise in transaction data. Feature engineering allowed us to extract meaningful insights, such as transaction time, location patterns, and unusual spending behaviors, which are critical in distinguishing fraud from legitimate transactions.

By employing machine learning models like Random Forest and XGBoost, we evaluated the performance of fraud detection techniques and found that advanced classification models significantly outperform traditional rule-based methods. Furthermore, balancing precision and recall was a key focus to minimizing false positives, ensuring a better customer experience while still effectively preventing fraud.

Additionally, real-time fraud detection was explored as a critical aspect of preventing financial loss before it occurs. By integrating streaming data analysis and lightweight predictive models, financial institutions can detect fraudulent activity as it happens and take proactive measures.

In conclusion, data science provides a scalable, adaptive, and intelligent solution to combat credit card fraud. With continuous advancements in machine learning and data analytics, financial institutions can enhance fraud detection accuracy, reduce financial losses, and improve consumer trust in digital transactions.